

*Each client scores both candidates on its own held-out data*

**Two candidate settings, both servable at the same cost**

shared head, no adapter

shared head, over the  
client's own adapter

**Client j**

run each candidate on the held-out examples,  
obtaining encrypted logits

select the logit of the true class with a plaintext mask  
the client knows its own labels, so the mask is public

compare against the encrypted maximum and  
accumulate the encrypted counts, one per class

encrypted  
counts

*No client ever sees a candidate model, and no per-example  
label is decrypted.*

*Combined and compared under encryption*

**Aggregation of the encrypted scores**

weight each class by its share of the  
federation's examples, taken from the  
public count matrix

the shared candidate pools evidence  
across clients; the personal candidate  
cannot, and its unseen classes score zero

one encrypted comparison of the two

one decrypted value:  
which candidate the federation adopts

*The two scores themselves are never revealed, so no client's  
accuracy on its own data is disclosed to the others.*