

HE-OFT: Privacy-Preserving One-Shot Federated Fine-Tuning under Homomorphic Encryption

Halil İbrahim Kanpak , Sinem Sav , Alptekin Küpçü 

Abstract—Organizations holding complementary data over the same task would each gain from a jointly fine-tuned model, yet cannot pool their data to build one. Federated learning offers collaboration but concentrates its privacy risk at training time. Gradients and per-round updates permit reconstruction of training data, and published attacks infer membership from them far more accurately than from the final model. A one-shot protocol that exchanges a single encrypted contribution removes that surface, but it still ends by handing the final model to every participant, which is not permitted where the model is itself a regulated or proprietary asset. Multi-round encrypted training also withholds the model, but it pays a cryptographic cost on every round and cannot adapt a pretrained backbone. We present HE-OFT, the first cryptographically secure one-shot federated fine-tuning protocol, in which the final model is never disclosed to any party. Each client fine-tunes a low-rank adapter and a classifier head on a frozen public backbone, retains the adapter locally, and uploads one encrypted head displacement. The server combines these under multiparty CKKS and never decrypts the result. Queries are answered under encryption, and a quorum of clients returns only the predicted label, addressed to the party that asked. Across four text classification tasks and one vision task, HE-OFT reaches 0.61 to 0.79 accuracy where a client training alone reaches 0.20 to 0.48, and it gives up 0.03 to 0.14 against a disclosed model. One query costs 31.5 s at four classes and 113.2 s at a hundred, and 13.5 MiB of traffic.

Index Terms—Federated learning, one-shot federated learning, federated fine-tuning, homomorphic encryption, multiparty computation, privacy-preserving machine learning.

I. INTRODUCTION

Organizations adapt large public pretrained models to their own data by fine-tuning [1], [2]. Often several organizations hold complementary data over the same task, different patient populations, fraud patterns, or customer bases, and each would obtain a better model from the combination than from its own data alone [3]. They cannot pool that data, because data protection law restricts the sharing of

confidential records across institutions [4], [5], [6], [7], [8]. Competitive concerns point the same way, and handing the data to whoever operates the fine-tuning infrastructure raises the same objection [9]. Three obstacles stand between that setting and a protocol these organizations could deploy.

Federated learning (FL) enables collaboration [10], [11], and its privacy risk is concentrated at training time [12],

[13]. The protocol exchanges model updates instead of data, but the updates are themselves the vulnerability. A gradient permits reconstruction of the batch that produced it [14], [15]. An observer of the per-round updates mounts membership inference at 87% accuracy on one model and dataset, where the same attack against the final model alone falls to 54.5% [12]. A server, free to choose the weights it broadcasts, can harvest client inputs outright [16], [17]. Two things therefore set the attack surface, namely what a protocol reveals while training runs and how often it reveals it.

Removing that surface takes two measures together. Making the protocol *one-shot*, so that each client contributes once, means that there is no intermediate state for anyone to observe. Encrypting that single contribution means that the one thing that is exchanged is never available in plaintext. In this paper, one-shot does not mean the parties stop communicating. It means that no intermediate training artifact is ever exposed. Prior work applies one of these measures at a time. One-shot FL [18], [19] sends its single contribution in plaintext [20], [21], [22], or protects it with differential privacy [23], [24], [25], which adds noise to the very artifact it releases. Encrypted FL runs many rounds of encrypted training or encrypted aggregation, paying the cryptographic cost on every round, and the direct route of training under encryption forces polynomial replacements of every activation and a deep, repeatedly bootstrapped circuit [26], [27], [28]. That cost scales with the network being trained. POSEIDON trains a three-layer network of sixty-four neurons per layer on handwritten digits at ten parties in 1.4 hours [26], and its successors lower the constant without changing what sets it. Adapting a pretrained transformer by this route is out of reach, and that is the case this paper addresses.

Applying both measures still leaves one artifact exposed, namely the model itself. One-shot protocols end by handing the final model to every participant, which is unacceptable when the model may not be distributed at all. Encrypted training does withhold it, but the cost just described puts a pretrained backbone out of reach. A model fine-tuned on confidential records can be a regulated artifact in high-risk domains such as health, biometrics, and credit [29], and it can be the asset the participants are least willing to give to each other. Such deployments need a protocol that builds a model and never releases it.

In this work, we present HE-OFT, the first one-shot federated fine-tuning protocol that enables several parties to adapt a pretrained model together and to query the result, without any of them ever holding it in plaintext. The parties fine-tune on a frozen public backbone, contribute once, and obtain a

Corresponding Author: Sinem Sav (sinem.sav@cs.bilkent.edu.tr)

Halil İbrahim Kanpak is with the Department of Computer Engineering, Koç University, İstanbul, Türkiye (e-mail: hkanpak21@ku.edu.tr).

Sinem Sav is with the Department of Computer Engineering, Bilkent University, Ankara, Türkiye (e-mail: sinem.sav@cs.bilkent.edu.tr).

Alptekin Küpçü is with the Department of Computer Engineering, Koç University, İstanbul, Türkiye (e-mail: akupcu@ku.edu.tr).

Manuscript received XX XX, XXXX; revised XX XX, XXXX.

shared classifier. The contribution is encrypted under multi-party CKKS [30], [31], a homomorphic encryption scheme whose secret key is split across the parties. The backbone stays public and in plaintext, so the cost of the cryptography does not grow with it. The classifier answers queries under encryption and returns a label, so a participant learns what the model predicts on its own inputs and nothing else about the model. There is no round in which a training artifact is exposed, no plaintext contribution, and no plaintext model at any point in the protocol.

Figure 1 shows how the shared head is built. Every client fine-tunes an adapter and a classifier head on the same frozen public backbone, keeps the adapter, and uploads one encrypted head displacement weighted by the classes it holds. The server adds the ciphertexts and divides by the per-class totals under encryption, so the shared head exists only as a ciphertext and the server holds no key to it. A client that wants a prediction computes the features of its query itself, encrypts them, and sends them to the serving party, which applies the encrypted head and takes the argmax under encryption. A quorum of clients then switches the result to that client alone, which decrypts one label.

The querier must build its query from parts it can run itself, so those parts must

be public. The shared trained quantity therefore sits after the last nonlinearity, which, in a transformer, is a single place. The three contributions below follow from that.

- **A split trainable unit.** Each client keeps its own adapter and never transmits it. Only the classifier head is federated, because a client can improve its own features alone but cannot recognize a class it has never seen. No adapter aggregate exists, so a coalition of all but one client cannot recover the last client's adapter by subtracting its own contributions (section III-B).
- **A serving path that never decrypts.** The querying client encrypts its own features, so the serving party cannot read the query. The serving party applies the encrypted head and takes the argmax under encryption, and a quorum switches the result to the querier's key. Only a label leaves the protocol, so a client cannot collect the per-class scores that would let it solve for the head (section III-C).
- **A selection rule that runs under encryption.** Two arrangements are servable at the same cost, the shared head over each client's own adapter and the shared head over the bare public backbone. Which one is better depends on the task. Each client scoring both on its own held-out data and voting estimates the wrong quantity, and selects wrongly whenever the two quantities differ. The estimator given here corrects that, needs no fitted threshold, and reveals only which arrangement won (section III-D). We measure all three in real multiparty CKKS, including the traffic every query costs (section V-E).

We evaluate on four text classification tasks and one vision task, with frozen RoBERTa [32] and vision transformer (ViT) [33] backbones, and we implement the cryptographic protocol in real multiparty CKKS rather than simulating it. We report what withholding the model costs, in accuracy against

TABLE I
NOTATION.

Symbol	Meaning
N	number of participating clients
$\mathcal{D}_j, n_j$	client j 's dataset and its size $n_j = \mathcal{D}_j $
C	number of classes
$n_{j,c}$	examples of class c held by client j
φ	frozen public backbone, never trained or encrypted
A_j	client j 's adapter, trained locally and never transmitted
φ_j	client j 's feature map, the backbone composed with A_j
θ_0	public initialiser of the classifier head
h_j	client j 's trained head
Δ_j	head displacement $h_j - \theta_0$
w_j	sample weight $n_j / \sum_i n_i$
θ^*	shared head, held only as a ciphertext
$\text{Enc}(\cdot)$	encryption under the collective public key

a disclosed model and in the communication and computation it adds, including the traffic that recurs on every query.

II. PRELIMINARIES

A. Notation

Table I collects the symbols. Client indices are $j \in \{1, \dots, N\}$ and class indices are $c \in \{1, \dots, C\}$.

B. Multiparty Homomorphic Encryption

We build on CKKS [30], a ring-learning-with-errors construction [34] that performs approximate arithmetic on vectors of real numbers under encryption. CKKS packs many real slots into one ciphertext and supports addition of ciphertexts, multiplication of a ciphertext by a plaintext, and multiplication of two ciphertexts. It is a levelled scheme, so each multiplication consumes one level of a finite budget, and a depleted budget is restored by bootstrapping [35].

A single-key scheme would require one party to hold the secret key and therefore to be trusted with decryption. We use the multiparty variant of CKKS [31], where the secret key is shared among the clients through a distributed key-generation protocol that produces a collective public key without ever assembling the secret in one place. Any party, the server included, may compute on ciphertexts under the collective public key, but decryption requires the cooperation of a quorum of clients and is carried out as a collective key-switching protocol. Two further protocols of the same family are used below. The first is a key switch to a designated public key, which re-encrypts a result so that one chosen party can read it. The second is a collective refresh, which restores a depleted level budget and plays the role bootstrapping plays in the single-key setting.

The construction is stated for a t -out-of- N threshold CKKS scheme with $2 \leq t \leq N$. Encryption is under a single collective public key. Decryption requires partial decryption shares from at least t clients, and any set of at most $t - 1$ clients learns nothing about a plaintext. Key generation, key switching and collective refresh are the threshold protocols of the same scheme.

Two consequences follow, and the protocol uses both. Any t clients form a quorum, so serving does not require every




Fig. 1. Building the shared head. Every client fine-tunes on the same frozen public backbone. The adapter is trained locally and never transmitted, so no aggregate of the adapters is ever formed and a coalition of clients has no sum from which to subtract its own contributions. Only the head displacement is encrypted and uploaded. The server adds the ciphertexts, weighting each by a public count, and divides by the per-class totals under encryption, holding no decryption key. The result is a shared head that no party decrypts.

client to be online. Any $t - 1$ clients are powerless against a ciphertext, so $t - 1$ is the corruption bound throughout section IV.

Our implementation instantiates $t = N$, which is the setting the multiparty CKKS library provides, and the measurements of section V-E are taken there. Every statement in section IV holds for general t .

III. METHOD

The protocol runs in two phases. In the first, each client fine-tunes an adapter and a classifier head on the same frozen public backbone, keeps the adapter, and uploads its head displacement once, weighted by its own per-class counts and encrypted. The server adds the ciphertexts and divides by the per-class totals under encryption, so the shared head exists only as a ciphertext (fig. 1). In the second, a client computes the features of its query itself, encrypts them, and sends them to the serving party, which applies the encrypted head and reduces the encrypted logits to the index of the largest (fig. 2). A quorum of clients switches that index to the querying client, which decrypts one label. The federation also chooses between two servable arrangements under encryption, and decrypts only which one it adopts (fig. 3). The subsections below derive each step.

A. Setting

We consider N clients holding private labeled data over a common label space of C classes. Client j holds data $\mathcal{D}_j$ of

size n_j . The class proportions differ across clients, and a client may hold no example of a class. The clients want a classifier that works across the whole label space, and they cannot pool their data to build one.

Three requirements separate this setting from ordinary federated learning. (1) The protocol is *one-shot*: each client uploads once, and there is no iterative exchange. (2) Confidentiality is *cryptographic*. A client's contribution is never available in plaintext to any other party, and we do not consider differential privacy due to the accuracy loss incurred. (3) The resulting model is *never disclosed*: no party, client or server, ever holds the trained classifier in plaintext.

Each requirement forecloses a family of designs, and the remaining space is narrow. We record the chain of implications here, so that the choices made in the rest of this section can be read as consequences of the requirements.

- C1** *Optimization cannot run under encryption at acceptable cost.* Gradient steps, softmax, and adaptive curvature require deep, repeatedly bootstrapped homomorphic circuits. *Therefore*, for efficiency, all learning happens in plaintext on the clients, and the server performs no learning.
- C2** *A server computing on ciphertexts cannot read its inputs.* It cannot inspect, compare, or branch, so it cannot select a data-dependent aggregation rule. *Therefore*, the choice is deferred to the clients, who make it without decrypting anything (section III-D).
- C3** *Models trained from different starting points do not com-*

bine linearly. In the parameter space of a deep network they occupy incomparable regions. *Therefore*, every client starts its trainable unit from one public initializer on one frozen public backbone, and the linear combination of **C4** applies only after that common start.

- C4** *Multiplicative depth dominates homomorphic cost.* *Therefore*, the aggregation is a linear combination, and the design is arranged so that this cheapest operation is exactly the one the method needs.
- C5** *Intermediate training signals are the strongest attack surface.* Each exposed round is a fresh training-time artifact, far more revealing than a finished model (section VI-A). *Therefore*, the protocol is one-shot, and its single artifact is a ciphertext.
- C6** *No single party may be able to decrypt.* The parties are mutually distrustful and the server is semi-honest. *Therefore*, the secret key is generated and held distributively, under multiparty CKKS (section II-B).
- C7** *The model is never disclosed, yet it must answer queries.* A query is formed from features, and the client that asks the query must be able to compute those features itself. *Therefore*, the client runs the backbone and its own adapter in plaintext, and the trained quantity that is shared cannot sit inside the backbone.

This requirement distinguishes the present design from one in which the model is released, and section III-B derives its consequences.

The three contributions of this paper follow from these constraints, a partition of the trainable unit (section III-B), a procedure for answering queries without decrypting the model (section III-C), and a procedure by which the federation selects between candidate arrangements it cannot observe (section III-D).

B. Partitioning the Trainable Unit

a) *The construction:* Each client uses the same frozen public backbone

φ , and trains two things on its own data, a low-rank adapter A_j that adjusts the representation, and a classifier head h_j on top of it. Only the head is shared. Client j encrypts the displacement $\Delta_j = h_j - \theta_0$ of its head from the public initializer and uploads it once. The server combines the encrypted displacements into a single shared head and never decrypts the result. The adapter is trained locally and never transmitted. In fig. 1 the adapters stay inside the client boxes, only the head displacement crosses to the server, and the arrow that would carry an adapter aggregate does not exist.

b) *Necessity of the partition:* The client that sends an inference query must compute the features of that query itself, by **C7**. If the shared trained quantity were an adapter placed inside the backbone, the features would depend on it, and computing them would require evaluating the backbone with encrypted weights. Every nonlinearity of the network would then have to be evaluated homomorphically, which is the cost this design exists to avoid. The trained quantity that is shared must therefore attach after the last nonlinearity, and in a transformer, the only such point is the final linear map. That is the head.

The requirement constrains the *position* of the shared map, not the task it serves. Any trained linear map that sits after the last nonlinearity satisfies it, provided the querying client can compute the features that enter the map. A classifier head is one such map. The vocabulary projection of an autoregressive decoder is another, since it is linear, sits after the final layer normalisation, and takes features the client can compute from a prefix it generated itself. We evaluate classification only. The construction itself does not depend on the label space being a set of classes, and section V-H records what a generation setting would additionally require, including the restriction to greedy decoding that section III-C imposes.

The adapter is relocated. Each client retains its own, so the representation still adapts to local data, and the federation shares the part where sharing is necessary. The distinction follows from coverage. A client can improve its own representation alone, but it cannot learn to recognise a class it has never observed. Coverage is a property of the head, which is exactly what the federation supplies. Section V-B measures both halves, how far a client gets with its own representation and how much the shared head adds on the classes it does not hold.

Retaining the adapter locally has a second consequence, which is cryptographic. No aggregate of the adapters is ever formed. A coalition of $N - 1$ clients therefore has no adapter sum from which to subtract its own contributions, and cannot recover the remaining client's adapter by arithmetic. Section IV states this precisely.

c) *The aggregation:* Let $g_j \in \mathbb{R}^C$ collect client j 's per-class example counts, $g_{j,c} = n_{j,c}$. The shared head is the coverage-weighted combination

$$\theta^* = \theta_0 + \frac{\sum_{j=1}^N g_j \odot \Delta_j}{\sum_{j=1}^N g_j}, \quad (1)$$

in which row c of the shared head is the count-weighted average of the clients' row- c displacements. Clients holding a class decide its row, and rows that no client covers remain at the public initializer. Each client forms the products $g_j \odot \Delta_j$ locally, before encryption, so the server only adds ciphertexts.

d) *The division:* The denominator of eq. (1) is the vector of federation-wide per-class totals. It cannot be decrypted. Every client knows its own counts, so a coalition of $N - 1$ clients that learned the totals would subtract its own and read the remaining client's class histogram exactly, which for an institutional participant is a description of its case mix. The reciprocal is therefore evaluated under encryption, using an inversion circuit over the positive domain whose refreshes are collective. This is the one deep operation in the construction of the shared head, and it is paid once, at aggregation, never per query. Section V-E reports its measured cost.

e) *Cost:* The server performs one ciphertext addition per client per ciphertext, one encrypted reciprocal, and one multiplication. The encrypted object is the head alone, whose size is set by the number of classes and the feature dimension and not by the backbone behind it. Communication is one upload per client and no download of model parameters at all, since no client ever receives the shared head.

Algorithm 1 Building the shared head. All clients hold the same frozen public backbone φ and the same public head initialiser θ_0 .

-
- 1: clients run distributed key generation, obtaining a collective public key
 - 2: **for** client $j = 1, \dots, N$ **in parallel do**
 - 3: train adapter A_j and head h_j on $\mathcal{D}_j$ ▷ plaintext, local
 - 4: **keep** A_j , which is never transmitted
 - 5: send $\text{Enc}(g_j \odot \Delta_j)$ and $\text{Enc}(g_j)$, where $\Delta_j = h_j - \theta_0$
 - 6: **end for**
 - 7: **server** adds the ciphertexts to form the numerator and the denominator of eq. (1) ▷ additions only
 - 8: **server** evaluates the reciprocal of the denominator under encryption and multiplies ▷ once, not per query
 - 9: **return** $\text{Enc}(\theta^*)$, which no party decrypts
-

C. Inference on the Encrypted Head

a) *The construction:* A client that wants a prediction computes the features of its query with its own backbone and adapter, encrypts them, and sends the ciphertext to a serving party. The serving party applies the encrypted shared head to the encrypted features, obtains encrypted logits, and reduces them to the index of the largest entry under encryption. A quorum of clients then key-switches that index to the querying client, which decrypts one label. Figure 2 shows the path, and algorithm 2 states it.

b) *Encryption of the query:* The serving party could be given the features in plaintext, which would make the head application cheaper. We encrypt them because features are invertible. A party holding $\varphi_j(x)$ and the public backbone can reconstruct an approximation of x . Sending features in plaintext would therefore hand the serving party the client's input, which is the thing the protocol exists to protect. The price is that the head is applied ciphertext by ciphertext rather than ciphertext by plaintext, and section V-E reports what that costs.

c) *Restriction to the predicted label:* Returning logits would defeat the construction. The head is a linear map on features the client computes itself, so a client that reads logits for feature vectors of its choosing recovers the head by solving a linear system in as many queries as the feature dimension [36]. This is not a hypothetical concern. The same object, the final projection layer of a production transformer, has been recovered through a deployed interface for a few tens of dollars, using an interface that exposed log-probabilities and a logit bias, and the operators' response was to restrict that combination [37]. Returning only the index of the largest entry is the same restriction imposed by construction rather than by policy. It places an adversary in the hard-label setting, where functionally equivalent extraction is markedly harder and has only recently been achieved for general networks [38], [39]. It does not make extraction impossible, and section IV states the resulting bound honestly, as a bound on queries rather than a cryptographic hardness claim.

d) *Directed decryption:* Decryption needs a quorum, so other clients necessarily participate in producing an answer to a question they did not ask. Rather than decrypt the label collectively, which would reveal it to every participant, the

Algorithm 2 Answering one query. The shared head remains a ciphertext throughout.

-
- Require:** query x at client j ; $\text{Enc}(\theta^*)$ held by the serving party
- 1: client j computes $\varphi_j(x)$ with its own backbone and adapter ▷ plaintext, local
 - 2: client j sends $\text{Enc}(\varphi_j(x))$
 - 3: **serving party** computes $\text{Enc}(\ell) \leftarrow \text{Enc}(\theta^*) \text{Enc}(\varphi_j(x))$ ▷ ciphertext by ciphertext
 - 4: **serving party** reduces $\text{Enc}(\ell)$ to $\text{Enc}(\hat{y})$, $\hat{y} = \arg \max_c \ell_c$, by a tournament of homomorphic comparisons [40]
 - 5: a quorum of clients key-switches $\text{Enc}(\hat{y})$ to client j 's public key
 - 6: **return** $\hat{y}$, readable by client j alone
-

quorum key-switches the ciphertext to the querying client's public key. The other clients contribute shares and learn nothing.

e) *Cost and the role of the serving party:* The argmax is the one operation in the protocol that is not shallow. Evaluated as a sequential fold of $C - 1$ comparisons it costs minutes for a large label space. Packing the logits into one ciphertext and reducing them by a tournament lowers the number of sequential comparisons from $C - 1$ to $\lceil \log_2 C \rceil$, and section V-E reports the measured result. None of this is particular to the multiparty setting. Homomorphic evaluation under a collectively generated key is identical to evaluation under a single key, and the schemes differ only in key generation and in decryption, so single-key results carry over. The serving party is not trusted with anything. It holds the collective public key, the evaluation keys, and ciphertexts, so it computes but cannot decrypt, and it learns neither the model nor any client's data. It is trusted only for availability, and where honest computation must be certified, verifiable homomorphic encryption applies [41].

f) *Scope:* This construction serves a trained quantity that is linear in the features. Serving a trained quantity that sits inside the backbone is the problem that secure transformer inference addresses, and an encrypted shared adapter composes with any such system at that literature's cost. We do not re-solve it. One difference should be noted. Those systems evaluate a plaintext model on encrypted inputs, whereas here the shared weights are encrypted as well, so each adapter site would carry a ciphertext-by-ciphertext product. The dominant cost in that literature is the nonlinearities, which are unchanged, so the added cost is the low-rank product per site rather than a different order of magnitude.

D. Candidate Selection Under Encryption

a) *The construction:* Two arrangements are servable at the same cost, the shared head over each client's own adapter, and the shared head over the bare public backbone with no adapter at all. Which is better depends on the task, and the federation must choose without decrypting either. Each client evaluates both on data it holds back from its own training, entirely under encryption, and reports encrypted per-class counts. The counts are combined into one encrypted score per arrangement, the two scores are compared under encryption, and exactly one value is decrypted, which arrangement the




Fig. 2. Answering one query. The querying client computes the features with its own backbone and adapter and encrypts them, so the serving party never sees the feature vector and cannot invert it to recover the query. The serving party holds the evaluation keys only: it applies the encrypted head, obtains encrypted logits, and reduces them to the index of the largest entry by a tournament of homomorphic comparisons. A quorum of clients then switches that index to the querying client's key, so the other clients contribute shares without learning the answer.

federation adopts. Figure 3 shows this, and algorithm 3 states it.

b) Inadequacy of the held-out vote: The natural procedure is for each client to measure the accuracy of both arrangements on its held-out data and vote. The procedure estimates the wrong quantity. The arrangement without an adapter is a single model, and the union of the clients' data is the whole training set, so pooled held-out measurements estimate its accuracy on the global distribution closely. The arrangement with personal adapters is N different models, and client j 's held-out data can only measure client j 's own model. How client j 's model performs on the other clients' distributions is not observable locally, and that is precisely where it fails. A client's held-out set also contains only the classes that client holds, so it cannot measure performance on classes the client has never seen. The measurement is systematically optimistic for the personal arrangement, and section V-D shows that the resulting vote chooses wrongly whenever the two differ.

c) A prior-weighted estimator: Score both arrangements as the expected accuracy of a prediction on an example drawn from the federation's own class distribution,

$$E = \sum_{c=1}^C p_c a_c, \quad p_c = \frac{\sum_j n_{j,c}}{\sum_j n_j}, \quad (2)$$

where a_c is the accuracy on class c and p_c is that class's share of the federation's examples, taken from the per-class counts already used in eq. (1). Two asymmetries follow from the setting. For the shared arrangement there is one model, so per-class evidence from different clients concerns the same object and combines, and every class is covered by some client. For

Algorithm 3 Choosing an arrangement without decrypting either. Nothing is revealed except the index of the winner.

```

1: for each arrangement  $m$  and each client  $j$  do
2:   for each held-out example  $(x, y)$  of client  $j$  do
3:     obtain  $\text{Enc}(\ell)$  and  $\text{Enc}(\max_c \ell_c)$  as in algorithm 2
4:     select  $\text{Enc}(\ell_y)$  with a plaintext mask  $\triangleright$  the client knows  $y$ , so the mask is public
5:     add  $\text{Enc}(1[\ell_y = \max_c \ell_c])$  to the class- $y$  accumulator
6:   end for
7: end for
8: combine the accumulators into  $\text{Enc}(E_m)$  by eq. (2), using public scalars  $\triangleright$  depth one
9: compare the two encrypted scores
10: return the decrypted index of the larger  $\triangleright$  one value

```

the personal arrangement each client's evidence concerns only its own model, and classes it does not hold contribute no evidence at all. Scoring those classes at zero makes E a lower bound for the personal arrangement, so the federation adopts it only when even a pessimistic estimate exceeds the accurately measured alternative. There is no fitted threshold anywhere. The rule compares two numbers on the same absolute scale.

d) Requirements on the held-out split: The estimator needs an estimate of accuracy on classes a client has not seen, and its only local proxy is accuracy on classes it has barely seen. A held-out set carved uniformly almost never contains a class the client holds one or two examples of, so that proxy is unmeasurable. Each client therefore reserves at least one held-out example from every class it holds before carving the remainder at random. A client holding a single example of a class donates it, so its adapter never trains on that class, which is exactly the probe the estimator needs. This costs a




Fig. 3. Choosing between two arrangements without decrypting either. Each client scores both on data held back from its own training, selecting the logit of the true class with a plaintext mask, which is public because the client knows its own labels. The per-class counts are weighted by each class's share of the federation's examples, taken from the count matrix already used for the aggregation. The shared arrangement pools evidence across clients because it is one model; the personal arrangement cannot, and classes a client does not hold contribute nothing, which makes its score a lower bound. Exactly one value is decrypted.

negligible amount of training data and requires no change to the cryptography.

e) Cost and disclosure: Every step is depth one except the comparisons, which are the same circuit already used for the argmax. The client knows its own labels, so the mask that selects the true class is a public plaintext, and the per-class totals are the same aggregate already used in eq. (1), so the weighting is by public scalars. One value is decrypted for the whole federation. The per-client per-class accuracies are not decrypted, and must not be. Together with the count matrix they would describe each client's label distribution alongside its model quality.

The threat model, the view each party has and what a coalition of clients can learn, is stated in section IV.

IV. SECURITY

The participants are N clients and a server. The server is semi-honest. It follows the protocol and may try to infer private information from what it observes. Clients may likewise be semi-honest and may collude. Section IV-C strengthens the client side of this model, and allows up to $t - 1$ clients to deviate arbitrarily while the server stays honest. Section IV-D explains why that assumption cannot be dropped.

a) View of the server: Ciphertexts under the collective public key, the public per-client sample counts, and the evaluation keys. It holds no decryption key, and by section II-B it

cannot decrypt. The server additionally sees encrypted query features, which it cannot read and therefore cannot invert.

The design is arranged so that no arithmetic shortcut evades it. No aggregate of the adapters exists, so there is no sum from which a coalition can subtract its own contributions. The shared head is never decrypted, so it cannot be inverted the same way. The per-class totals are never decrypted, so a coalition cannot recover the remaining client's class histogram.

A. The Ideal Functionality

Functionality 1 defines the ideal functionality $\mathcal{F}$ against which the protocol is measured. The server and the serving party are entities outside $\mathcal{F}$, so either may be corrupt. The form of the functionality follows the secure aggregation functionality of ELSA [42], and placing the aggregator outside it follows FULLSA [43].

The sample counts sit in the leakage because the protocol uses them as public scalars. The phase signals sit there because anyone watching the network sees them. We prove security relative to both.

Three properties of $\mathcal{F}$ are what the protocol must reproduce. No party ever holds θ^* . The only value the whole federation learns is one index. A client learns the labels it asks for, up to its allowance, and nothing else.

Functionality 1 HE-OFT $\mathcal{F}$

Parameters. The client count N , the decryption threshold t , the per-client query allowance Q , and the label space size C . The parties are clients $P_1, \dots, P_N$, an aggregation server S and a serving party V , none of which is part of $\mathcal{F}$. The functionality holds a counter q_j for each client, each initialised to 0, and a store that is initially empty and that $\mathcal{F}$ never sends to any party.
Inputs. Each client P_j inputs its dataset $\mathcal{D}_j$, and later inputs queries x . The server S inputs a selection request. The serving party V inputs nothing.
Outputs. Every party receives the selected index a^* and the phase signals. Client P_j receives one label for each of its first Q queries, and $\perp$ afterwards.
Steps.

- 1) *Training and aggregation.* On $(\text{Train}, \mathcal{D}_j)$ from every client, run the local training of algorithm 1 on each $\mathcal{D}_j$ and store the shared head θ^* of eq. (1). Send $(\text{Done}, \text{train})$ to every party.
- 2) *Selection.* On (Select) from S , evaluate the estimator of eq. (2) for both servable arrangements, send the index a^* of the larger to every party, and send $(\text{Done}, \text{select})$ to every party.
- 3) *Serving.* On (Query, x) from client P_j , return $\perp$ to P_j if $q_j \geq Q$. Otherwise set $q_j \leftarrow q_j + 1$, return $\hat{y} = \arg \max_c (\theta^* \varphi_j(x))_c$ to P_j alone, and send $(\text{Done}, \text{query}, j)$ to V .

Leakage. $\mathcal{F}$ reveals $\mathcal{L}$ to the adversary, consisting of the public parameters, the per-client sample counts n_j , the announced index a^* , and the fact that each Done message was sent.

B. Security Against Semi-Honest Parties

Definition 1 (Semi-honest security). Let $\mathcal{A}$ corrupt a set of parties that may contain the server, the serving party, and at most $t - 1$ clients. Write $\text{view}_{\mathcal{A}}^{\Pi}(\{\mathcal{D}_j\}_{j=1}^N)$ for the corrupted parties' joint view of a protocol run on datasets $\{\mathcal{D}_j\}_{j=1}^N$. The protocol Π securely realizes $\mathcal{F}$ with leakage $\mathcal{L}$ if there is a probabilistic polynomial-time simulator $\mathcal{S}$ such that, for every $\{\mathcal{D}_j\}_{j=1}^N$,

$$\text{view}_{\mathcal{A}}^{\Pi}(\{\mathcal{D}_j\}_{j=1}^N) \stackrel{c}{\approx} \mathcal{S}(\mathcal{L}, \{\mathcal{D}_j\}_j \text{ corrupt}, \text{out}_{\mathcal{A}}),$$

where $\text{out}_{\mathcal{A}}$ collects the answers $\mathcal{F}$ returns to the corrupted clients.

Theorem 1 (Semi-honest security). *Assume the t -out-of- N threshold CKKS scheme is secure under chosen plaintext attack (IND-CPA) against an adversary holding $t - 1$ key shares. Then the protocol securely realizes $\mathcal{F}$ with leakage $\mathcal{L}$ against any semi-honest adversary corrupting the server, the serving party, and at most $t - 1$ clients, in the sense of definition 1.*

Proof sketch. The simulator is given $\mathcal{L}$, the corrupted clients' datasets, and the answers the corrupted clients receive. It proceeds in three parts.

Setup. The simulator runs the distributed key generation with the corrupted parties, playing the honest clients on uniformly chosen shares. That transcript is simulatable by the security of the key-generation protocol.

Training. For each honest client the simulator sends an encryption of zero in place of $\text{Enc}(g_j \odot \Delta_j)$ and $\text{Enc}(g_j)$. A hybrid argument replaces the honest ciphertexts one at a time, and each neighbouring pair of hybrids is indistinguishable by IND-CPA against $t - 1$ shares, the assumption of the theorem.

Selection and serving. The selection step decrypts one value. The simulator takes that index from $\mathcal{L}$ and produces the collective key-switching transcript for it from that protocol's simulator. For each query of a corrupted client, the simulator holds the answer in $\text{out}_{\mathcal{A}}$ and simulates the key switch to that

client in the same way. Queries by honest clients are switched to keys the adversary does not hold, so their transcripts are simulated as key switches of encryptions of zero.

Composing the three parts gives a view computationally indistinguishable from the real one. $\square$

C. Input Privacy Against Malicious Clients

Theorem 1 assumes every party follows the protocol. We now let the clients deviate, and section IV-D shows why the server must stay honest. We bound what a deviating coalition learns about the data of a client that does not deviate. We do not claim correctness.

Definition 2 (Content game). The adversary $\mathcal{A}$ corrupts $t - 1$ clients. The server and the serving party are honest. At least one client h stays honest.

- 1) $\mathcal{A}$ outputs two datasets D_0 and D_1 with $|D_0| = |D_1|$.
- 2) The challenger samples $b \in \{0, 1\}$ and gives D_b to client h .
- 3) The protocol runs. $\mathcal{A}$ may deviate arbitrarily in the values its clients upload, in the shares they contribute, and in the queries they ask, up to the allowance of every client it controls.
- 4) $\mathcal{A}$ outputs b' . Its advantage is $|\Pr[b' = b] - \frac{1}{2}|$.

Theorem 2 (Input privacy). *Assume the conditions of theorem 1. Then the advantage of $\mathcal{A}$ in definition 2 is at most $\text{negl}(\lambda) + \delta(Q_{\text{tot}})$, where $Q_{\text{tot}} = (t - 1)Q$ and δ is the advantage available from Q_{tot} label queries to the served model.*

Proof sketch. The adversary's view consists of the key-generation transcript, the honest clients' uploaded ciphertexts, the values the server computes, the key-switching transcripts, and the answers to its own queries.

The first four are independent of b up to $\text{negl}(\lambda)$.

What remains is the answers. These do depend on b , because θ^* depends on client h 's displacement. That is the dependence of any prediction interface on its training data, and it is bounded by what Q_{tot} label queries reveal, which is $\delta(Q_{\text{tot}})$ by definition. $\square$

The bound scales with the corruption threshold. A coalition of $t - 1$ clients commands $(t - 1)Q$ queries whatever the size of the federation. Section V-F measures δ and turns it into a bound on Q .

D. Necessity of an Honest Server

Theorem 2 keeps the server honest. That assumption is necessary, and the reason is worth stating, because it marks the boundary of what this message pattern can achieve.

Decryption is a protocol the clients run on a ciphertext the server presents. A deviating server can present a ciphertext of its own choosing, such as a row of θ^* , rather than the output of algorithm 2. The natural repair is to have an honest client refuse such a request. It cannot, and the obstacle is the encryption itself.

Proposition 1 (No plaintext gate from a ciphertext). *Let an honest client decide, by a polynomial-time predicate D on its view, whether to contribute its decryption share for a presented ciphertext. Let the rest of that view be independent of the underlying plaintext. If the scheme is IND-CPA against an adversary holding $t - 1$ key shares, then for every pair of plaintexts m_0 and m_1 ,*

$$|\Pr[D(\text{Enc}(m_0)) = 1] - \Pr[D(\text{Enc}(m_1)) = 1]| \leq \text{negl}(\lambda).$$

Proof. A gap between the two probabilities gives an IND-CPA distinguisher that holds one key share. One share is at most $t - 1$ shares, so the assumed security of the scheme applies directly. $\square$

An honest client therefore releases its share for both plaintexts or for neither. No protocol of this message pattern realizes $\mathcal{F}$ against a malicious server, which is why theorem 2 places the server among the honest parties.

The serving circuit is a deterministic function of the encrypted head, the encrypted query, the evaluation keys and the public parameters, and its level restoration is deterministic as well, because the protocol restores levels under collectively generated bootstrapping keys rather than by collective refresh. A client can therefore recompute the prescribed output ciphertext and compare it as a string.

E. Leakage Inherent to the Functionality

$\mathcal{F}$ answers queries. Answers carry information about θ^* , and enough of them reconstruct it. Section V-F measures the rate at three to five queries per parameter for a functional copy. This is a property of the functionality, not of our instantiation of it. Any protocol that realizes $\mathcal{F}$ inherits it.

The cryptographic claim is that the protocol adds nothing to the leakage of $\mathcal{F}$. The operational claim is that $\mathcal{F}$'s own leakage is bounded by the query allowance Q . The first rests on IND-CPA. The second rests on a measurement and on a deployment's ability to meter queries.

V. EXPERIMENTS

A. Setup

The trainable unit is a rank-8 low-rank adapter with its down-projection frozen at a shared public initialisation, together with a classifier head. No labelled public data is used anywhere in this paper.

We evaluate on four text classification tasks of increasing label-space size, AG-News with 4 classes, TREC with 6, DBpedia with 14 and Banking77 with 77, using a frozen RoBERTa-base encoder [32], and on CIFAR-100 with a frozen ViT-B/16 [33]. Data are partitioned across clients by a Dirichlet distribution over labels with concentration α , the standard label-skew partition [44], [45]. A smaller α leaves each client a few dominant classes and almost none of the rest. Unless stated otherwise there are $N = 10$ clients, $\alpha = 0.1$, the local trajectory is 200 steps, and every number is the mean over three seeds.

We write A_{sel} for the accuracy of the servable arrangement that section V-D selects. Three reference points appear

TABLE II

ACCURACY ON THE GLOBAL TEST SET, THREE SEEDS, $N = 10$, $\alpha = 0.1$. BOTH SERVABLE ARRANGEMENTS ARE SHOWN, TOGETHER WITH WHAT A CLIENT OBTAINS ALONE, WHAT A DISCLOSED MODEL WOULD REACH, AND WHAT ONE MODEL TRAINED ON THE POOLED DATA REACHES AT THE SAME TOTAL NUMBER OF STEPS. NEITHER REFERENCE IS SERVABLE UNDER THE THREAT MODEL OF SECTION IV. WE DID NOT RUN THE POOLED REFERENCE ON CIFAR-100. ON ONE SEED THE DIRICHLET DRAW ASSIGNS FEWER THAN TWENTY SAMPLES TO SOME CLIENTS, WHICH CANNOT TRAIN, SO THE EFFECTIVE FEDERATION IS SEVEN ON AG-NEWS AND NINE ON TREC.

Task	C	servable			reference	
		shared head	adapter	alone	disclosed	pooled
AG-News	4	0.649	0.479	0.475	0.739	0.921
TREC	6	0.607	0.429	0.400	0.711	0.967
DBpedia	14	0.789	0.754	0.451	0.925	0.988
Banking77	77	0.206	0.686	0.249	0.760	0.923
CIFAR-100	100	0.748	0.774	0.197	0.784	—

throughout, and none of them is servable under the threat model of section IV. We write A_{loc} for the accuracy a client obtains alone, with its own adapter and its own head and no federation. We write A_{dis} for the accuracy of the disclosed model, in which the adapter and the head are both aggregated and decrypted, so that every party holds the result in plaintext. We write A_{pool} for the accuracy of one model trained on the union of the clients' data at the same total number of steps.

Two differences follow, and the paper reports both. The quantity $A_{\text{dis}} - A_{\text{sel}}$ is the price of never disclosing a model. The quantity $A_{\text{pool}} - A_{\text{dis}}$ is the price of the partition.

B. Utility of the Federated Head

Sharing only the classifier head, over representations that each client adapts locally, yields a model usable across the whole label space, and the preferable arrangement depends on the size of the label space.

Table II reports both servable arrangements and both reference points. The arrangement labelled *shared head* gives every client the same model over the bare public backbone. The arrangement labelled *personal adapter* gives each client the shared head over its own locally adapted representation.

Three observations follow. First, the federation raises accuracy on every task. A client alone reaches 0.20 to 0.48, while the federated head reaches 0.61 to 0.79 on four of the five tasks. The gap is largest exactly where a client's own data covers least of the label space.

Second, the preferable arrangement changes with the size of the label space. On the small label spaces the shared head wins, 0.17 on AG-News and 0.18 on TREC. On the large ones the personal adapter wins. The shared head collapses to 0.206 while the personal adapter holds 0.686. The collapse is a coverage failure. With 77 classes and this skew each head row is decided by few clients or none. A locally adapted representation improves separability for every class the client will ever be asked about, including those it never observed.

Third, the two failure modes are complementary. The shared head fails when coverage is thin. The personal adapter fails when each client's representation drifts far enough that a common head no longer transfers. Neither arrangement dominates

TABLE III

CIFAR-10 ON TWO PUBLISHED PARTITIONS, THREE SEEDS. THE COLUMNS ARE THE QUANTITIES DEFINED IN SECTION V-A, ALL MEASURED ON OUR PROTOCOL AT THE PARTITION EACH PRIOR PAPER PUBLISHED.

N	α	A_{loc}	A_{sel}	A_{dis}	$A_{\text{dis}} - A_{\text{sel}}$
5	0.10	0.384	0.948	0.962	0.013
5	0.30	0.569	0.950	0.964	0.014
20	0.04	0.212	0.948	0.943	-0.005
20	0.16	0.399	0.952	0.959	0.007

the other, so the federation must select between them, which is the subject of section V-D.

a) *Sensitivity*: Figure 4 varies the protocol axes around the default on DBpedia. Accuracy falls with skew, from 0.971 at $\alpha = 1.0$ to 0.789 at $\alpha = 0.1$. Longer local training helps, from 0.718 at 100 steps to 0.870 at 400. Accuracy rises with the size of the federation, from 0.803 at $N = 10$ to 0.882 at $N = 50$.

C. Comparison with One-Shot Federated Learning

The protocol gives up no accuracy for protecting the contributions, where the differentially private one-shot methods give up an amount that grows as their budget tightens. The peer group evaluates on different tasks from ours, so we ran our protocol on theirs. We adopt two published partitions of CIFAR-10: $N = 5$ with $\alpha \in \{0.1, 0.3\}$, the configuration of [20], and $N = 20$ with $\alpha \in \{0.04, 0.16\}$, the configuration of [23]. Table III reports the quantities of section V-A on each, as the mean over three seeds.

At $N = 5$ and $\alpha = 0.1$ we measure $A_{\text{sel}} = 0.948$, against the 0.503 of [20] and the 0.406 that the same paper reports for [46]. We do not present that margin as a result. Those methods train a ResNet-18 from scratch, whereas every client here fine-tunes a frozen public ViT-B/16, and the representation accounts for essentially all of the difference. Absolute accuracy across papers therefore measures the backbone.

What compares across papers is the accuracy each method gives up for protecting the contribution, since every paper measures

that against its own unprotected baseline on its own architecture, and the backbone cancels. Read that way the peer group falls into three cases. A method either leaves the contribution unprotected, protects it with differential privacy, or protects it cryptographically.

FedAUXfdp does protect its contribution, and its cost is governed by the budget. Against its own undefended baseline of 0.752 on CIFAR-10 at $N = 20$ and $\alpha = 0.16$, the published figures give up 0.006 at $\epsilon = 1.0$ and 0.029 at $\epsilon = 0.5$, then 0.413 at $\epsilon = 0.1$ and 0.626 at $\epsilon = 0.01$ [23].

Our contribution is protected cryptographically, and that protection costs no accuracy, by construction, since the decrypted result matches the plaintext computation to the encoding precision of the scheme. Withholding the model is a separate charge, and no method above incurs it. At $N = 20$ and $\alpha = 0.16$ that charge is $A_{\text{dis}} - A_{\text{sel}} = 0.007$, against the 0.006 to 0.626 that the same paper's mechanism costs as ϵ

falls from 1.0 to 0.01. Over table III the charge lies between -0.005 and 0.014 , and over table II between 0.03 and 0.14 .

One difference bounds the federation's own contribution independently of the backbone, because both of its terms are measured on the same backbone. At $N = 5$ and $\alpha = 0.1$ we measure $A_{\text{loc}} = 0.384$ against $A_{\text{sel}} = 0.948$. That difference is what the protocol supplies.

D. Accuracy of the Selection Rule

Scoring both arrangements under the federation's own class prior, letting the shared arrangement pool evidence across clients and scoring the personal arrangement's unobserved classes at zero, selects correctly in 12 of 15 cells. Filling the unobserved classes from each client's rarest classes instead of with zero selects correctly in 13.

E. Cryptographic Cost

The cryptographic cost of the protocol is dominated by one-time setup, and the recurring per-query traffic is a few megabytes. We do not simulate the cryptography. The protocol is implemented in real multiparty CKKS on Lattigo [47], [31]. All figures use ring degree 2^{14} for the shallow operations and a deeper chain at ring degree 2^{15} for the circuits that need it. Timings are single-run wall clock on a commodity CPU and are reported as indicative. Communication figures are exact.

a) *What one query costs, end to end*: On one core a query takes 31.5 s at four classes and 113.2 s at a hundred. The arithmetic and the key switch account for 0.24 s of that at every label-space size. Everything else is the encrypted argmax, which is therefore the only operation whose cost decides whether a deployment is practical. The tournament costs about 16 s per round and needs $\lceil \log_2 C \rceil$ rounds, against the fourteenfold larger sequential fold (fig. 4).

b) *Communication*: The protocol is one-shot in the sense of section I: no intermediate training artifact is ever exposed. It is not silent. Key generation precedes it, selection costs a bounded exchange, and serving costs traffic per query.

The per-query total is $8.5 + 0.5N$ MiB, since the quorum returns one key-switching share each. It is 11.0 MiB at $N = 5$ and 18.5 MiB at $N = 20$, and it is 6.8 MiB at ring degree 2^{14} and 27.0 MiB at 2^{16} . It does not vary with the number of classes. Setup costs 344 MiB per client, once.

c) *Correctness*: The decrypted result matches the plaintext computation to a relative error at the level of the encoding precision of the scheme, and the encrypted argmax is exact. The accuracy figures of table II are therefore unaffected by the cryptography.

F. Residual Leakage

The protocol exposes no training-time artifact and no model, so the only remaining channel is the sequence of answers the served model returns, and that channel is bounded by the query allowance rather than eliminated.

Reaching fidelity 0.90 takes about 1.2×10^4 queries on AG-News, 5.0×10^4 on DBpedia and 2.0×10^5 on Banking77. The head has Cd parameters, and in each case fidelity 0.90




Fig. 4. The protocol along four axes, one record behind each panel. Panels (a) to (c) vary the federation size, the label skew and the length of the local trajectory on DBpedia around the default of $N = 10$, $\alpha = 0.1$ and 200 steps. Panel (a) and panel (c) are one seed and panel (b) is the mean over three. Panel (d) gives the cost of one encrypted argmax against the size of the label space, at $N = 10$ and ring degree 2^{15} , for the tournament the protocol uses and for the sequential fold it replaces.

costs between three and five times that number of queries. A deployment can therefore set its allowance from C and d without repeating this experiment.

a) *Comparison with a released model:* A participant holding the model in plaintext has fidelity 1 at zero queries, may read the parameters directly rather than infer them, and is subject to no allowance because there is nothing to meter.

G. Comparison with Model Disclosure

Never disclosing the model costs accuracy, and table II allows the price to be read directly. The quantity $A_{\text{dis}} - A_{\text{sel}}$ is 0.071 on AG-News, 0.104 on TREC, 0.136 on DBpedia, 0.075 on Banking77 and 0.029 on CIFAR-100. After the protocol ends, there is no plaintext model anywhere.

H. Scope and Limitations

a) *Linearity of the shared quantity:* This is what section III-B shows the never-disclosed requirement forces, and it is a real restriction. A shared adapter inside the backbone would adapt the representation for every participant rather than only for its owner. Serving such an adapter reduces to secure transformer inference (section VI-E) and composes with that literature at its cost.

b) *Bounds on model extraction:* The query allowance of section V-F is an operational control.

c) *Malicious participants:* The cryptographic guarantees assume semi-honest parties. A client that uploads a crafted head displacement can bias the shared head, and because contributions are encrypted the server cannot inspect them. Two directions are compatible with the encryption and are left to future work, an upload-time norm bound enforced by a zero-knowledge proof, and robust aggregation evaluated homomorphically.

VI. RELATED WORK

A. Attacks at Training Time and at Inference Time

Federated learning (FL) began as an iterative procedure in which a server repeatedly broadcasts a model, clients train it on their own data, and the server averages the returned updates [10]. In every round each client sends a new update computed on its private data, and several attacks have shown

that these updates leak more than anything else the protocol reveals. A gradient allows reconstruction of the batch that produced it [14], [15]. An observer of the per-round updates infers membership and unintended properties of another client's data, and the repetition of those updates is what makes the attack strong [12], [13], [48]. A server that chooses the weights it broadcasts does not need inference at all, since crafted weights make clients return verbatim copies of their own inputs, even through a secure-aggregation layer [16], [17]. Several works share predictions instead of gradients, and the predictions leak too [49], [50], [51], [52].

Attacks that see only the final model are consistently weaker, and the field survey draws that separation explicitly [11]. Several works infer membership from a released model, or invert it from released confidences [53], [54], [55], [56], [57], and against models that generalise well they identify few members, and only at very low false-positive rates [58]. The reason for the gap is simple. An update is a detailed view of one client's data and arrives every round, where the final model is a single average that generalisation moves away from any one training example. A protocol that never sends an update, and that sends anything only once, therefore leaves every adversary with the weaker attack.

B. Protecting Training with Cryptography

Secure aggregation lets the server learn only the sum of the clients' updates, using masks that cancel in the aggregate [59]. It is lossless, but it protects one round's sum inside an iterative protocol, and the revealed per-round sums have themselves been shown to leak across rounds [60], [61]. Homomorphic encryption goes further by computing on contributions that stay encrypted. POSEIDON runs encrypted stochastic gradient descent under multiparty CKKS, with activations replaced by polynomial approximations so that they can be evaluated homomorphically [26]. The guarantee is strong and the model loses no accuracy, but the approach is iterative and encrypted nonlinearity costs a deep multiplicative circuit with repeated bootstrapping. Several works measure how delicate those polynomial activations are [62], [63], [64], [65], [66]. That cost scales with the network being trained, so adapting a pretrained transformer by this route is out of reach.

Several works encrypt only the aggregation step of an otherwise plaintext training loop [67], [68], [69], [70], [71]. All of them remain bound to the iterative protocol, so the encryption cost is paid on every round and what is protected is a per-round sum rather than a complete contribution. One of them also shows why key structure matters, since single-key homomorphic encryption forces the assumption that the server and the clients do not collude [70]. The threshold structure of section II-B removes that assumption.

Transfer learning itself has been placed under encryption. Several works train a classifier head, a linear or logistic model, or a low-rank adapter on frozen features under CKKS [72], [73], [74], [75], [76], [77], [78], [79]. These share our structural niche, a frozen public backbone with a small trainable unit, but they run an optimiser on ciphertexts, so multiplicative depth grows with the step count and every nonlinearity must be approximated. Several of them are also single-party, one data owner outsourcing computation rather than a collaboration of mutually distrustful parties. Our parties train on their own and must protect what they trained.

C. One-Shot Federated Learning and Differential Privacy

A parallel line collapses the many rounds into a single exchange [18], [19]. Several works exchange predictions on a shared public set, which in practice repeats the exchange [80], [81], [46]. Genuinely single-round methods distil the clients' models through a generator or synthesised inputs, or fuse the models directly, in one round or in stages the authors account as one [20], [21], [82], [83], [22]. These establish that one round can produce a usable model, but they work in plaintext, and each client sends an entire model or a synthetic distillate of its data. By the evidence of section VI-A that is what an adversary would target. A recent method sends no model at all, having each client run a frozen pretrained encoder and upload per-class sufficient statistics once [84]. Those statistics travel in plaintext, and per-class totals are exactly what section III-B keeps encrypted.

The one-shot methods that do address privacy use differential privacy. FedAUXfdp releases client contributions under a differentially private mechanism after distillation through a pretrained extractor [23], and related approaches protect a diffusion-generated surrogate or a teacher ensemble [25], [24], or apply a noise-free mechanism to the distillation itself [85]. These are our natural quantitative peers, and their privacy is lossy, because the noise a meaningful budget requires is added to the very artifact that is released. Differential privacy proves a statistical bound on how much a released artifact can reveal and buys that bound with accuracy. Encryption makes a contribution computationally indistinguishable from random to every party without the key and costs the model nothing. Several works apply the same trade-off to transfer learning, running noisy gradient descent on frozen features or on a low-rank adapter [86], [87], [88], [89], [90], [91]. They share our backbone-and-adapter structure, but they pay a budget-dependent accuracy tax on the model they release, and they are centralised or multi-round.

D. Federated Fine-Tuning and Task Arithmetic

A recent line federates parameter-efficient fine-tuning, communicating a low-rank adapter rather than a model. FedIT averages the adapters round by round [92], but per-factor averaging is biased, because a client learns a product of two factors and averaging the factors separately does not give the average of the products. Several works correct that bias by stacking factors, by reconstructing full products, by reconciling heterogeneous ranks, or by sharing one factor only [93], [94], [95], [96]. FFA-LoRA instead freezes the randomly initialised down-projection and trains only the up-projection, so that averaging the trained factor is exact [97]. Combining displacements from a shared initialiser is task-vector merging, which is equivalent to one-shot federated averaging, and one round of fine-tuning suffices for foundation models in plaintext [98], [99], [100]. None of this line protects the contributions, which is our subject.

E. Secure Inference

Answering queries from a model that stays encrypted is the subject of secure inference. Several works evaluate a transformer under encryption, either non-interactively under CKKS or between two or three parties with secret sharing, and in each case the nonlinearities of attention dominate the cost [101], [102], [103], [104]. This literature is what makes the disclosure setting of section III-C affordable, and we compose with it. The map served here is linear in the features, so serving a trained map that sits inside the backbone instead would reduce to exactly the problem these systems address.

Two systems sit closer. slytHERin evaluates a network under CKKS with the model held under a multiparty key, and it key-switches the prediction to the querier, which is the serving setting of section III-C [105]. The same group built the encrypted training scheme of section VI-B, so the pair is the closest existing route to a model that no party decrypts. Two properties separate it from what follows. It evaluates the whole network homomorphically, which is why the models it reports are a twenty-layer and a fifty-layer network on a handwritten-digit task rather than a pretrained encoder. And it returns the prediction vector, so the querier receives scores, which section V-F shows is a materially cheaper target for extraction than a label is.

CryptPEFT divides the work as we do and reaches the opposite conclusion about who holds the secret. A public frozen backbone runs in plaintext on the client, the client encrypts the resulting features, only the adapter is evaluated under encryption, and the client receives the predicted label alone [106]. Confining encrypted computation to a small unit after the representation is the motivation for our design as well. It is a two-party protocol, the trained unit is the provider's plaintext, and there is no training protocol, so it establishes that the serving arrangement is practical without addressing how mutually distrustful parties would build the unit at all.

Constructions for the encrypted maximum faster than the tournament of section III-C exist, both for ranking and order statistics at constant comparison depth and for decoding a language model without decrypting the scores [107], [108]. Both

TABLE IV

THE CLOSEST SYSTEMS ON THE FOUR AXES THIS PAPER TURNS ON, AS EACH PAPER REPORTS THEM. MODEL IN THE CLEAR SAYS WHETHER ANY PARTY ENDS UP HOLDING THE TRAINED MODEL IN PLAINTEXT, AND *provider* MEANS ONE PARTY DOES WHILE THE QUERIER DOES NOT. SLYtHERIN AND CRYPTPEFT SERVE A MODEL THEY DO NOT TRAIN, SO THEIR ROUNDS CELL IS EMPTY.

System	Rounds	Protection	Model in the clear	Querier receives
DENSE [20]	one	none	yes	the model
GH-OFL [84]	one	none	yes	the model
FedAUXfdp [23]	one	DP	yes	the model
POSEIDON [26]	many	HE	yes	the model
slytHERin [105]		HE	no	scores
CryptPEFT [106]		2PC	provider	label
HE-OFT	one	HE	no	label

are single-key and neither carries the collective refreshes that the threshold setting requires, so neither transfers unchanged.

Positioning. Table IV places the closest systems on the axes that separate them. The multi-round cryptographic line protects every round but pays on every round, and what it protects is a per-round sum. The one-shot line exposes its single contribution in plaintext, or noises it and pays in accuracy. The fine-tuning line communicates a small object but protects nothing. Every one of them ends by handing the finished model to the participants, so the model itself remains a disclosed artifact. A single prior scheme is at once one-shot, federated and encrypted, but only for a single-layer closed-form learner that cannot adapt a deep representation [109]. What remains unaddressed is a protocol that exposes no training-time quantity, that exchanges anything at all only once, and that never discloses the model even to the parties that built it. Section III works out what that forces, and the rest of the paper measures what it costs.

VII. CONCLUSION

We presented HE-OFT, a federated fine-tuning protocol in which the trained model is never disclosed. The design follows from the observation that federated learning’s privacy risk is concentrated in the quantities a protocol exposes while the model is being built, and from the further requirement that in some deployments the finished model may not be distributed either. Each client fine-tunes on a frozen public backbone, retains its adapter locally, and uploads a single encrypted head displacement. The server combines those displacements under multiparty CKKS and never decrypts the result. Queries are answered under

encryption, and a quorum of clients returns only the predicted label, addressed to the party that asked.

Withholding the model constrains what may be shared. Since the querying party must form its query from components it can evaluate itself, those components must be public, and the shared trained quantity must therefore attach after the last nonlinearity of the network. This is the source of the design’s three parts, the partition of the trainable unit, inference on the encrypted head, and selection between candidate arrangements that no party can observe. Across four text classification

tasks and one vision task, the protocol produces a model substantially stronger than any client obtains alone, at a cost of 0.03 to 0.14 accuracy relative to a model that is aggregated and disclosed. The cryptographic layer itself costs no accuracy, and the recurring communication is a few megabytes per query.

Two directions follow. Robustness to actively malicious participants is outside the present guarantee, and reconciling robust aggregation with an encrypted contribution remains open. An upload-time norm bound enforced by a zero-knowledge

proof is one route. Serving a shared quantity that sits inside the backbone, rather than after it, reduces to secure transformer inference, and the accuracy that would recover is quantified in table II.

ACKNOWLEDGMENTS

We acknowledge the support of TÜBİTAK (the Scientific and Technological Research Council of Türkiye) projects 124E091 and 124N941. The authors used Claude Sonnet 5 and Claude Opus 5 to enhance the manuscript by shortening sentences, correcting typos, and improving grammar.

REFERENCES

- [1] R. Bommasani, D. A. Hudson, E. Adeli *et al.*, “On the opportunities and risks of foundation models,” *arXiv preprint arXiv:2108.07258*, 2021.
- [2] E. J. Hu, Y. Shen, P. Wallis, Z. Allen-Zhu, Y. Li, S. Wang, L. Wang, and W. Chen, “LoRA: Low-rank adaptation of large language models,” in *ICLR*, 2022.
- [3] M. J. Sheller, B. Edwards, G. A. Reina, J. Martin, S. Pati, A. Kotrotsou, M. Milchenko, W. Xu, D. Marcus, R. R. Colen, and S. Bakas, “Federated learning in medicine: facilitating multi-institutional collaborations without sharing patient data,” *Scientific Reports*, vol. 10, p. 12598, 2020.
- [4] European Parliament and Council of the European Union, “Regulation (eu) 2016/679 (general data protection regulation),” *Official Journal of the European Union*, L 119, pp. 1–88, 2016.
- [5] U.S. Department of Health and Human Services, “HIPAA privacy rule,” 45 C.F.R. Parts 160 and 164, 2013.
- [6] Republic of Türkiye, “Law on the protection of personal data (kişisel verilerin korunması kanunu), law no. 6698,” 2016.
- [7] Standing Committee of the National People’s Congress, “Personal information protection law of the people’s republic of china,” 2021.
- [8] Parliament of India, “Digital personal data protection act,” 2023.
- [9] G. Xiao, J. Lin, and S. Han, “Offsite-tuning: Transfer learning without full model,” *arXiv preprint arXiv:2302.04870*, 2023.
- [10] B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y. Arcas, “Communication-efficient learning of deep networks from decentralized data,” in *AISTATS*, 2017.
- [11] P. Kairouz, H. B. McMahan, B. Avent *et al.*, “Advances and open problems in federated learning,” *Foundations and Trends in Machine Learning*, vol. 14, no. 1–2, pp. 1–210, 2021.
- [12] M. Nasr, R. Shokri, and A. Houmansadr, “Comprehensive privacy analysis of deep learning: Passive and active white-box inference attacks against centralized and federated learning,” in *IEEE S&P*, 2019.
- [13] L. Melis, C. Song, E. De Cristofaro, and V. Shmatikov, “Exploiting unintended feature leakage in collaborative learning,” in *IEEE S&P*, 2019.
- [14] L. Zhu, Z. Liu, and S. Han, “Deep leakage from gradients,” in *NeurIPS*, 2019.
- [15] J. Geiping, H. Bauermeister, H. Dröge, and M. Moeller, “Inverting gradients – how easy is it to break privacy in federated learning?” in *NeurIPS*, 2020.
- [16] F. Boenisch, A. Dziedzic, R. Schuster, A. S. Shamsabadi, I. Shumailov, and N. Papernot, “When the curious abandon honesty: Federated learning is not private,” in *IEEE EuroS&P*, 2023.
- [17] L. Fowl, J. Geiping, W. Czaja, M. Goldblum, and T. Goldstein, “Robbing the fed: Directly obtaining private data in federated learning with modified models,” in *ICLR*, 2022.

- [18] N. Guha, A. Talwalkar, and V. Smith, "One-shot federated learning," *arXiv preprint arXiv:1902.11175*, 2019.
- [19] F. Amato, L. Qiu, M. Tanveer, S. Cuomo, D. Annunziata, F. Giampaolo, and F. Piccialli, "Towards one-shot federated learning: Advances, challenges, and future directions," *Neurocomputing*, vol. 664, p. 132088, 2026.
- [20] J. Zhang, C. Chen, B. Li, L. Lyu, S. Wu, S. Ding, C. Shen, and C. Qi, "DENSE: Data-free one-shot federated learning," in *NeurIPS*, 2022.
- [21] R. Dai, Y. Zhang, A. Li, T. Liu, X. Yang, and B. Han, "Enhancing one-shot federated learning through data and ensemble co-boosting," in *ICLR*, 2024.
- [22] Z. Tang, Y. Zhang, P. Dong, Y.-m. Cheung, A. C. Zhou, B. Han, and X. Chu, "Fusefl: One-shot federated learning through the lens of causality with progressive model fusion," in *NeurIPS*, 2024.
- [23] H. Hoech, R. Rischke, K. Mueller, and W. Samek, "FedAUXfdp: Differentially private one-shot federated distillation," in *IJCAI Workshop on Federated Learning*, 2022.
- [24] Q. Li, B. He, and D. Song, "Practical one-shot federated learning for cross-silo setting," in *IJCAI*, 2021.
- [25] M. Mendieta, G. Sun, and C. Chen, "Navigating heterogeneity and privacy in one-shot federated learning with diffusion models," in *WACV*, 2025.
- [26] S. Sav, A. Pyrgelis, J. R. Troncoso-Pastoriza, D. Froelicher, J.-P. Bossuat, J. Sá Sousa, and J.-P. Hubaux, "POSEIDON: Privacy-preserving federated neural network learning," in *NDSS*, 2021.
- [27] G. Xu, X. Han, S. Xu, T. Zhang, H. Li, X. Huang, and R. H. Deng, "Hercules: Boosting the performance of privacy-preserving federated learning," *IEEE TDS*, vol. 20, no. 5, pp. 4418–4433, 2023.
- [28] A. Pirillo and L. Colombo, "ReBoot: Encrypted training of deep neural networks with CKKS bootstrapping," *arXiv:2506.19693*, 2025.
- [29] European Parliament and Council of the European Union, "Regulation (eu) 2024/1689 (artificial intelligence act)," *Official Journal of the European Union*, L, 2024/1689, 2024.
- [30] J. H. Cheon, A. Kim, M. Kim, and Y. Song, "Homomorphic encryption for arithmetic of approximate numbers," in *ASIACRYPT*, 2017.
- [31] C. Mouchet, J. R. Troncoso-Pastoriza, J.-P. Bossuat, and J.-P. Hubaux, "Multiparty homomorphic encryption from ring-learning-with-errors," *PoPETS*, vol. 2021, no. 4, pp. 291–311, 2021.
- [32] Y. Liu, M. Ott, N. Goyal, J. Du, M. Joshi, D. Chen, O. Levy, M. Lewis, L. Zettlemoyer, and V. Stoyanov, "RoBERTa: A robustly optimized BERT pretraining approach," *arXiv:1907.11692*, 2019.
- [33] A. Dosovitskiy, L. Beyer, A. Kolesnikov, D. Weissenborn, X. Zhai, T. Unterthiner, M. Dehghani, M. Minderer, G. Heigold, S. Gelly, J. Uszkoreit, and N. Houlsby, "An image is worth 16x16 words: Transformers for image recognition at scale," in *ICLR*, 2021.
- [34] V. Lyubashevsky, C. Peikert, and O. Regev, "On ideal lattices and learning with errors over rings," *Journal of the ACM*, vol. 60, no. 6, pp. 43:1–43:35, 2013.
- [35] J. H. Cheon, K. Han, A. Kim, M. Kim, and Y. Song, "Bootstrapping for approximate homomorphic encryption," in *EUROCRYPT*, 2018.
- [36] F. Tramèr, F. Zhang, A. Juels, M. K. Reiter, and T. Ristenpart, "Stealing machine learning models via prediction APIs," in *USENIX Security*, 2016.
- [37] N. Carlini, D. Paleka, K. D. Dvijotham, T. Steinke, J. Hayase, A. F. Cooper, K. Lee, M. Jagielski, M. Nasr, A. Conmy, I. Yona, E. Wallace, D. Rolnick, and F. Tramèr, "Stealing part of a production language model," in *ICML*, 2024.
- [38] Y. Chen, X. Dong, J. Guo, Y. Shen, A. Wang, and X. Wang, "Hard-label cryptanalytic extraction of neural network models," in *ASIACRYPT*, 2024.
- [39] N. Carlini, J. Chávez-Saab, A. Hamblitz, F. Rodríguez-Henríquez, and A. Shamir, "Polynomial time cryptanalytic extraction of deep neural networks in the hard-label setting," in *EUROCRYPT*, 2025.
- [40] J. H. Cheon, D. Kim, and D. Kim, "Efficient homomorphic comparison methods with optimal complexity," in *ASIACRYPT*, 2020.
- [41] A. Viand, C. Knabenhans, and A. Hithnawi, "SoK: Fully homomorphic encryption compilers and verifiable computation," in *IEEE S&P*, 2023.
- [42] M. Rathee, C. Shen, S. Wagh, and R. A. Popa, "ELSA: Secure aggregation for federated learning with malicious actors," in *IEEE S&P*, 2023.
- [43] F. Karakoç, A. Küpçü, and M. Önen, "Fault tolerant and malicious secure federated learning," in *CANS*, 2024.
- [44] T.-M. H. Hsu, H. Qi, and M. Brown, "Measuring the effects of non-identical data distribution for federated visual classification," *arXiv preprint arXiv:1909.06335*, 2019.
- [45] Q. Li, Y. Diao, Q. Chen, and B. He, "Federated learning on non-IID data silos: An experimental study," in *IEEE ICDE*, 2022.
- [46] T. Lin, L. Kong, S. U. Stich, and M. Jaggi, "Ensemble distillation for robust model fusion in federated learning," in *NeurIPS*, 2020.
- [47] C. Mouchet, J.-P. Bossuat, J. Troncoso-Pastoriza, and J.-P. Hubaux, "Lattigo v5: A multiparty homomorphic encryption library in Go," in *WAHC*, 2021.
- [48] O. Zari, C. Xu, and G. Neglia, "Efficient passive membership inference attack in federated learning," in *PriML Workshop, NeurIPS*, 2021.
- [49] H. Takahashi, J. Liu, and Y. Liu, "Breaching FedMD: Image recovery via paired-logits inversion attack," in *CVPR*, 2023.
- [50] Z. Yang, Y. Zhao, and J. Zhang, "FD-Leaks: Membership inference attacks against federated distillation learning," in *APWeb-WAIM*, 2023.
- [51] Y. Gu and Y. Bai, "LDIA: Label distribution inference attack against federated learning in edge computing," *Journal of Information Security and Applications*, vol. 74, p. 103475, 2023.
- [52] H. Shi, T. Ouyang, and A. Wang, "Unveiling client privacy leakage from public dataset usage in federated distillation," *PoPETS*, vol. 2025, no. 4, pp. 201–215, 2025.
- [53] R. Shokri, M. Stronati, C. Song, and V. Shmatikov, "Membership inference attacks against machine learning models," in *IEEE S&P*, 2017.
- [54] S. Yeom, I. Giacomelli, M. Fredrikson, and S. Jha, "Privacy risk in machine learning: Analyzing the connection to overfitting," in *IEEE CSF*, 2018.
- [55] A. Salem, Y. Zhang, M. Humbert, P. Berrang, M. Fritz, and M. Backes, "ML-Leaks: Model and data independent membership inference attacks and defenses on machine learning models," in *NDSS*, 2019.
- [56] M. Fredrikson, S. Jha, and T. Ristenpart, "Model inversion attacks that exploit confidence information and basic countermeasures," in *ACM CCS*, 2015.
- [57] B. Hitaj, G. Ateniese, and F. Pérez-Cruz, "Deep models under the GAN: Information leakage from collaborative deep learning," in *ACM CCS*, 2017.
- [58] N. Carlini, S. Chien, M. Nasr, S. Song, A. Terzis, and F. Tramèr, "Membership inference attacks from first principles," in *IEEE S&P*, 2022.
- [59] K. Bonawitz, V. Ivanov, B. Kreuter, A. Marcedone, H. B. McMahan, S. Patel, D. Ramage, A. Segal, and K. Seth, "Practical secure aggregation for privacy-preserving machine learning," in *ACM CCS*, 2017.
- [60] J. So, R. E. Ali, B. Güler, J. Jiao, and A. S. Avestimehr, "Securing secure aggregation: Mitigating multi-round privacy leakage in federated learning," in *AAAI*, 2023.
- [61] R. Kerkouche, G. Ács, and M. Fritz, "Client-specific property inference against secure aggregation in federated learning," in *WPES*, 2023.
- [62] K. Garimella, N. K. Jha, and B. Reagen, "Sisyphus: A cautionary tale of using low-degree polynomial activations in privacy-preserving deep learning," in *PPML Workshop, ACM CCS*, 2021.
- [63] M. Baruch, N. Drucker, L. Greenberg, and G. Moshkovich, "A methodology for training homomorphic encryption friendly neural networks," in *ACNS Workshops*, 2022.
- [64] P. Agamennone, "Polynomial approximations of relu for secure cnn inference in homomorphic encryption environments," *IEICE Trans. Fundamentals*, vol. E108.A, no. 7, pp. 977–980, 2025.
- [65] F. Al Hossain and T. Rahman, "A training framework for optimal and stable training of polynomial neural networks," *arXiv preprint arXiv:2505.11589*, 2025.
- [66] A. Ibarrondo and M. Önen, "FHE-compatible batch normalization for privacy-preserving deep learning," in *PriML Workshop, NeurIPS*, 2021.
- [67] C. Zhang, S. Li, J. Xia, W. Wang, F. Yan, and Y. Liu, "Batchcrypt: Efficient homomorphic encryption for cross-silo federated learning," in *USENIX ATC*, 2020.
- [68] W. Jin, Y. Yao, S. Han, J. Gu, C. Joe-Wong, S. Ravi, S. Avestimehr, and C. He, "Fedml-he: An efficient homomorphic-encryption-based privacy-preserving federated learning system," in *NeurIPS*, 2023.
- [69] X. Wei, R. Huang, L. Lei, and J. Wu, "FedSHE-CQ: A communication-efficient homomorphic encryption framework for federated learning," *Computer Networks*, vol. 274, p. 111813, 2026.
- [70] J. Liu, L. Yan, B. Li, L. Yu, and C. Shen, "SHE-LoRA: Selective homomorphic encryption for federated tuning with heterogeneous LoRA," in *ICLR*, 2026.
- [71] I. Emmaka and Tran Viet Xuan Phuong, "One-shot secure aggregation: A hybrid cryptographic protocol for private federated learning in IoT," in *IEEE TPS-ISA*, 2025.
- [72] S. Lee, G. Lee, J. W. Kim, J. Shin, and M.-K. Lee, "HETAL: Efficient privacy-preserving transfer learning with homomorphic encryption," in *ICML*, 2023.

- [73] M. Kim, Y. Song, S. Wang, Y. Xia, and X. Jiang, "Secure logistic regression based on homomorphic encryption: Design and evaluation," *JMIR Medical Informatics*, vol. 6, no. 2, p. e19, 2018.
- [74] A. Kim, Y. Song, M. Kim, K. Lee, and J. H. Cheon, "Logistic regression model training based on the approximate homomorphic encryption," *BMC Medical Genomics*, vol. 11, no. Suppl 4, p. 83, 2018.
- [75] J. Chiang, "Privacy-preserving CNN training with transfer learning: Two hidden layers," *arXiv preprint arXiv:2504.12623*, 2025.
- [76] A. M. I. M. Osmani, T. Rahman, M. S. Rahman, and S. Islam, "PrivFedTL: Privacy-preserving federated transfer learning for resource constrained devices," *Computer Networks*, vol. 286, p. 112449, 2026.
- [77] Al Amin, K. Hasan, L. Hong, and S. Ullah, "Privacy-preserving federated vision transformer learning leveraging lightweight homomorphic encryption in medical AI," in *IEEE ICNC*, 2026.
- [78] Y. Li, W. Yu, and J. Zhao, "PrivTuner with homomorphic encryption and loRA: A P3EFT scheme for privacy-preserving parameter-efficient fine-tuning of AI foundation models," *arXiv preprint arXiv:2410.00433*, 2024.
- [79] J. Frery, R. Bredehoft, J. Klemsa, A. Meyre, and A. Stoian, "Private loRA fine-tuning of open-source LLMs with homomorphic encryption," *arXiv preprint arXiv:2505.07329*, 2025.
- [80] D. Li and J. Wang, "FedMD: Heterogenous federated learning via model distillation," in *NeurIPS FL Workshop*, 2019.
- [81] S. Itahara, T. Nishio, Y. Koda, M. Morikura, and K. Yamamoto, "Distillation-based semi-supervised federated learning for communication-efficient collaborative training with non-IID private data," *IEEE Transactions on Mobile Computing*, vol. 22, no. 1, pp. 191–205, 2023.
- [82] X. Liu, L. Liu, F. Ye, Y. Shen, X. Li, L. Jiang, and J. Li, "FedLPA: One-shot federated learning with layer-wise posterior aggregation," in *NeurIPS*, 2024.
- [83] J. Zhang, S. Liu, and X. Wang, "One-shot federated learning via synthetic distiller-distillate communication," in *NeurIPS*, 2024.
- [84] F. Turazza, M. Picone, and M. Mamei, "The Gaussian-Head OFL family: One-shot federated learning from client global statistics," in *ICLR*, 2026.
- [85] L. Sun and L. Lyu, "Federated model distillation with noise-free differential privacy," in *IJCAI*, 2021.
- [86] F. Tramèr and D. Boneh, "Differentially private learning needs better features (or much more data)," in *ICLR*, 2021.
- [87] H. Mehta, W. Krichene, A. Thakurta, A. Kurakin, and A. Cutkosky, "Differentially private image classification from features," *TMLR*, 2023.
- [88] D. Yu, S. Naik, A. Backurs, S. Gopi, H. A. Inan, G. Kamath, J. Kulkarni, Y. T. Lee, A. Manoel, L. Wutschitz *et al.*, "Differentially private fine-tuning of language models," in *ICLR*, 2022.
- [89] X. Li, F. Tramèr, P. Liang, and T. Hashimoto, "Large language models can be strong differentially private learners," in *ICLR*, 2022.
- [90] X.-Y. Liu, R. Zhu, D. Zha, J. Gao, S. Zhong, M. White, and M. Qiu, "Differentially private low-rank adaptation of large language model using federated learning," *ACM TMIS*, vol. 16, no. 2, pp. 1–24, 2025.
- [91] J. Xu, K. Saravanan, R. van Dalen, H. Mehmood, D. Tuckey, and M. Ozy, "DP-DyLoRA: Fine-tuning transformer-based models on-device under differentially private federated learning using dynamic low-rank adaptation," in *IEEE ICC*, 2026.
- [92] J. Zhang, S. Vahidian, M. Kuo, C. Li, R. Zhang, T. Yu, Y. Zhou, G. Wang, and Y. Chen, "Towards building the federated GPT: Federated instruction tuning," in *IEEE ICASSP*, 2024.
- [93] Z. Wang, Z. Shen, Y. He, G. Sun, H. Wang, L. Lyu, and A. Li, "FLoRA: Federated fine-tuning large language models with heterogeneous low-rank adaptations," in *NeurIPS*, 2024.
- [94] J. Bai, D. Chen, B. Qian, L. Yao, and Y. Li, "Federated fine-tuning of large language models under heterogeneous tasks and client resources," in *NeurIPS*, 2024.
- [95] Y. J. Cho, L. Liu, Z. Xu, A. Fahrezi, and G. Joshi, "Heterogeneous LoRA for federated fine-tuning of on-device foundation models," in *EMNLP*, 2024.
- [96] P. Guo, S. Zeng, Y. Wang, H. Fan, F. Wang, and L. Qu, "Selective aggregation for low-rank adaptation in federated learning," in *ICLR*, 2025.
- [97] Y. Sun, Z. Li, Y. Li, and B. Ding, "Improving LoRA in privacy-preserving federated learning," in *ICLR*, 2024.
- [98] G. Ilharco, M. T. Ribeiro, M. Wortsman, L. Schmidt, H. Hajishirzi, and A. Farhadi, "Editing models with task arithmetic," in *ICLR*, 2023.
- [99] Z. Tao, I. Mason, S. Kulkarni, and X. Boix, "Task arithmetic through the lens of one-shot federated learning," *TMLR*, 2025.
- [100] Z. Wang, B. Tian, Y. He, Z. Shen, G. Sun, Y. Liu, L. Liu, M. Liu, and A. Li, "Revisiting federated fine-tuning: A single communication round is enough for foundation models," *arXiv preprint arXiv:2412.04650*, 2024.
- [101] J. Zhang, J. Liu, X. Yang, Y. Wang, K. Chen, X. Hou, K. Ren, and X. Yang, "Secure transformer inference made non-interactive," in *NDSS*, 2025.
- [102] M. Hao, H. Li, H. Chen, P. Xing, G. Xu, and T. Zhang, "Iron: Private inference on transformers," in *NeurIPS*, 2022.
- [103] Q. Pang, J. Zhu, H. Möllering, W. Zheng, and T. Schneider, "BOLT: Privacy-preserving, accurate and efficient inference for transformers," in *IEEE S&P*, 2024.
- [104] Y. Dong, W.-j. Lu, Y. Zheng, H. Wu, D. Zhao, J. Tan, Z. Huang, C. Hong, T. Wei, and W. Chen, "PUMA: Secure inference of LLaMA-7B in five minutes," *arXiv preprint arXiv:2307.12533*, 2023.
- [105] F. Intoci, S. Sav, A. Pyrgelis, J.-P. Bossuat, J. R. Troncoso-Pastoriza, and J.-P. Hubaux, "slytHERin: An agile framework for encrypted deep neural network inference," in *Cloud S&P*, 2023.
- [106] S. Xia, W. Wang, Z. Wang, Y. Zhang, Y. Jin, D. Meng, and R. Hou, "CryptPEFT: Efficient and private neural network inference via parameter-efficient fine-tuning," in *NDSS*, 2026.
- [107] F. Mazzone, M. Everts, F. Hahn, and A. Peter, "Efficient ranking, order statistics, and sorting under CKKS," in *USENIX Security*, 2025.
- [108] M. Avitan, M. Baruch, N. Drucker, I. Zimerman, and Y. Goldberg, "Efficient decoding methods for language models on encrypted data," *arXiv:2509.08383*, 2025.
- [109] O. Fontenla-Romero, B. Guijarro-Berdiñas, E. Hernández-Pereira, and B. Pérez-Sánchez, "FedHEONN: Federated and homomorphically encrypted learning method for one-layer neural networks," *Future Generation Computer Systems*, vol. 149, pp. 200–211, 2023.



HALİL İBRAHİM KANPAK received his B.Sc. degree in Mathematics from Koç University. He is currently pursuing his Ph.D. at Koç University, where he is a member of the Cryptography, Cyber Security & Privacy Research Group. His research interests include cryptography and related areas in Deep Learning.



Asst. Prof. SİNEM SAV received her Ph.D. in Computer and Communication Sciences from École Polytechnique Fédérale de Lausanne (EPFL). She also holds M.Sc. and B.Sc. degrees in Computer Engineering from Bilkent University, where she currently serves as an Assistant Professor in the Computer Engineering department. Her research interests include privacy-preserving machine learning and applied cryptography.



Prof. ALPTEKİN KÜPÇÜ received his Ph.D. degree from Brown University Computer Science Department in 2010. Since then, he has been working as a faculty member at Koç University, leading the Cryptography, Cyber Security & Privacy Research Group that he founded. He is an ACM and IEEE Senior Member, and a co-founder of Xtinge Technology Inc.